

act4.1

May 18, 2021

1 Actividad 4.1 - Redes neuronales artificiales

1.1 Edson Raul Cepeda Marquez 1820776

Construcción y entrenamiento de un modelo de red neuronal par la predicción de precios de casas utilizando el conjunto de datos de Boston_House_Prices.

1.2 Análisis de datos

1.2.1 Carga de datos

```
[4]: import pandas as pd
import numpy as np
from sklearn.datasets import load_boston
```

```
[3]: boston = load_boston()
```

```
[5]: data = pd.DataFrame(boston.data)
```

```
[6]: data.head()
```

```
[6]:
```

	0	1	2	3	4	5	6	7	8	9	10 \
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7

	11	12
0	396.90	4.98
1	396.90	9.14
2	392.83	4.03
3	394.63	2.94
4	396.90	5.33

Asigando precio y nombre de columnas

```
[7]: data.columns = boston.feature_names
```

```
[17]: data['PRICE'] = boston.target
```

```
[18]: data.head()
```

```
[18]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	\
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	

	PTRATIO	B	LSTAT	PRICE
0	15.3	396.90	4.98	24.0
1	17.8	396.90	9.14	21.6
2	17.8	392.83	4.03	34.7
3	18.7	394.63	2.94	33.4
4	18.7	396.90	5.33	36.2

```
[19]: print(data.shape)
```

(506, 14)

Corroborando si existen valores nulos

```
[20]: data.isnull().sum()
```

```
[20]:
```

CRIM	0
ZN	0
INDUS	0
CHAS	0
NOX	0
RM	0
AGE	0
DIS	0
RAD	0
TAX	0
PTRATIO	0
B	0
LSTAT	0
PRICE	0

dtype: int64

Estadísticas descriptivas

```
[21]: data.describe()
```

```
[21]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	\
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	
mean	3.613524	11.363636	11.136779	0.069170	0.554695	6.284634	

std	8.601545	23.322453	6.860353	0.253994	0.115878	0.702617
min	0.006320	0.000000	0.460000	0.000000	0.385000	3.561000
25%	0.082045	0.000000	5.190000	0.000000	0.449000	5.885500
50%	0.256510	0.000000	9.690000	0.000000	0.538000	6.208500
75%	3.677083	12.500000	18.100000	0.000000	0.624000	6.623500
max	88.976200	100.000000	27.740000	1.000000	0.871000	8.780000

	AGE	DIS	RAD	TAX	PTRATIO	B \
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000
mean	68.574901	3.795043	9.549407	408.237154	18.455534	356.674032
std	28.148861	2.105710	8.707259	168.537116	2.164946	91.294864
min	2.900000	1.129600	1.000000	187.000000	12.600000	0.320000
25%	45.025000	2.100175	4.000000	279.000000	17.400000	375.377500
50%	77.500000	3.207450	5.000000	330.000000	19.050000	391.440000
75%	94.075000	5.188425	24.000000	666.000000	20.200000	396.225000
max	100.000000	12.126500	24.000000	711.000000	22.000000	396.900000

	LSTAT	PRICE
count	506.000000	506.000000
mean	12.653063	22.532806
std	7.141062	9.197104
min	1.730000	5.000000
25%	6.950000	17.025000
50%	11.360000	21.200000
75%	16.955000	25.000000
max	37.970000	50.000000

```
[22]: data.info()
```


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 506 entries, 0 to 505
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   CRIM        506 non-null    float64
1   ZN          506 non-null    float64
2   INDUS       506 non-null    float64
3   CHAS        506 non-null    float64
4   NOX         506 non-null    float64
5   RM          506 non-null    float64
6   AGE         506 non-null    float64
7   DIS         506 non-null    float64
8   RAD         506 non-null    float64
9   TAX         506 non-null    float64
10  PTRATIO     506 non-null    float64
11  B           506 non-null    float64
12  LSTAT       506 non-null    float64
13  PRICE       506 non-null    float64
```

```
dtypes: float64(14)
memory usage: 55.5 KB
```

Observando la distribución de los precios


```
[26]: import seaborn as sns
      sns.displot(data=data.PRICE)
```

```
[26]: <seaborn.axisgrid.FacetGrid at 0x7f575af2e730>
```



```
[27]: sns.boxplot(data=data.PRICE)
```

```
[27]: <AxesSubplot:>
```



Calculando los coeficientes de correlación

```
[28]: correlation = data.corr()
```


```
[29]: correlation.loc['PRICE']
```

```
[29]: CRIM      -0.388305
      ZN        0.360445
      INDUS    -0.483725
      CHAS      0.175260
      NOX      -0.427321
      RM        0.695360
      AGE      -0.376955
      DIS       0.249929
      RAD      -0.381626
      TAX      -0.468536
      PTRATIO  -0.507787
      B         0.333461
      LSTAT    -0.737663
      PRICE     1.000000
      Name: PRICE, dtype: float64
```

```
[31]: import matplotlib.pyplot as plt
```


```
[33]: fig, axes = plt.subplots(figsize=(15,12))
      sns.heatmap(correlation, square=True, annot=True)
```

[33]: <AxesSubplot:>



Observando las variables más llamativas obtenidas del análisis de correlación

```
[34]: plt.figure(figsize=(20,5))
features = ['LSTAT', 'RM', 'PTRATIO']
for i, col in enumerate(features):
    plt.subplot(1, len(features), i+1)
    x = data[col]
    y = data.PRICE
    plt.scatter(x, y, marker='o')
    plt.title("Variation in House prices")
    plt.xlabel(col)
    plt.ylabel('"House prices in $1000"')
```



1.2.2 Preparando los datos para el entrenamiento

```
[36]: X = data.iloc[:, :-1]

[39]: y = data.PRICE

[40]: from sklearn.model_selection import train_test_split

[41]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2,
    ↪ random_state = 4)

[42]: from sklearn.preprocessing import StandardScaler

[43]: sc = StandardScaler()

[44]: X_train = sc.fit_transform(X_train)

[45]: X_test = sc.transform(X_test)
```

1.2.3 Construcción del modelo de red neuronal

```
[50]: import keras
    from keras.layers import Dense, Activation, Dropout
    from keras.models import Sequential

[51]: model = Sequential()
```

Definiendo la arquitectura de la red neuronal

```
[53]: model.add(Dense(128, activation='relu', input_dim = 13))
    model.add(Dense(64, activation='relu'))
    model.add(Dense(32, activation='relu'))
    model.add(Dense(16, activation='relu'))
    model.add(Dense(1))
    model.compile(optimizer='adam', loss='mean_squared_error')
```

Entrenando la red neuronal

```
[54]: model.fit(X_train, y_train, epochs=100)
```

```
Epoch 1/100
13/13 [=====] - 15s 2ms/step - loss: 586.9261
Epoch 2/100
13/13 [=====] - 0s 5ms/step - loss: 504.9662
Epoch 3/100
13/13 [=====] - 0s 3ms/step - loss: 329.5313
Epoch 4/100
13/13 [=====] - 0s 3ms/step - loss: 114.0957
Epoch 5/100
13/13 [=====] - 0s 3ms/step - loss: 55.6977
Epoch 6/100
13/13 [=====] - 0s 3ms/step - loss: 26.8154
Epoch 7/100
13/13 [=====] - 0s 2ms/step - loss: 17.7073
Epoch 8/100
13/13 [=====] - 0s 3ms/step - loss: 21.5348
Epoch 9/100
13/13 [=====] - 0s 3ms/step - loss: 16.5987
Epoch 10/100
13/13 [=====] - 0s 3ms/step - loss: 13.6926
Epoch 11/100
13/13 [=====] - 0s 3ms/step - loss: 11.2430
Epoch 12/100
13/13 [=====] - 0s 3ms/step - loss: 10.3592
Epoch 13/100
13/13 [=====] - 0s 3ms/step - loss: 11.7095
Epoch 14/100
13/13 [=====] - 0s 3ms/step - loss: 14.8494
Epoch 15/100
13/13 [=====] - 0s 3ms/step - loss: 11.6141
Epoch 16/100
13/13 [=====] - 0s 3ms/step - loss: 9.1067
Epoch 17/100
13/13 [=====] - 0s 3ms/step - loss: 9.6613
Epoch 18/100
13/13 [=====] - 0s 3ms/step - loss: 14.2203
Epoch 19/100
13/13 [=====] - 0s 3ms/step - loss: 9.2219
Epoch 20/100
13/13 [=====] - 0s 3ms/step - loss: 7.5127
Epoch 21/100
13/13 [=====] - 0s 3ms/step - loss: 11.1469
Epoch 22/100
13/13 [=====] - 0s 4ms/step - loss: 9.8665
```


Epoch 23/100
13/13 [=====] - 0s 3ms/step - loss: 9.1630
Epoch 24/100
13/13 [=====] - 0s 4ms/step - loss: 8.4459
Epoch 25/100
13/13 [=====] - 0s 3ms/step - loss: 8.4768
Epoch 26/100
13/13 [=====] - 0s 3ms/step - loss: 8.2416
Epoch 27/100
13/13 [=====] - 0s 3ms/step - loss: 8.9542
Epoch 28/100
13/13 [=====] - 0s 3ms/step - loss: 7.7439
Epoch 29/100
13/13 [=====] - 0s 3ms/step - loss: 8.6205
Epoch 30/100
13/13 [=====] - 0s 5ms/step - loss: 10.0104
Epoch 31/100
13/13 [=====] - 0s 3ms/step - loss: 6.8064
Epoch 32/100
13/13 [=====] - 0s 3ms/step - loss: 6.4711
Epoch 33/100
13/13 [=====] - 0s 4ms/step - loss: 7.3245
Epoch 34/100
13/13 [=====] - 0s 3ms/step - loss: 6.3654
Epoch 35/100
13/13 [=====] - 0s 3ms/step - loss: 8.3157
Epoch 36/100
13/13 [=====] - 0s 3ms/step - loss: 6.1892
Epoch 37/100
13/13 [=====] - 0s 3ms/step - loss: 7.7888
Epoch 38/100
13/13 [=====] - 0s 3ms/step - loss: 8.2177
Epoch 39/100
13/13 [=====] - 0s 4ms/step - loss: 6.7666
Epoch 40/100
13/13 [=====] - 0s 3ms/step - loss: 6.2496
Epoch 41/100
13/13 [=====] - 0s 3ms/step - loss: 6.5795
Epoch 42/100
13/13 [=====] - 0s 3ms/step - loss: 5.3427
Epoch 43/100
13/13 [=====] - 0s 3ms/step - loss: 6.4896
Epoch 44/100
13/13 [=====] - 0s 3ms/step - loss: 5.1552
Epoch 45/100
13/13 [=====] - 0s 3ms/step - loss: 6.7148
Epoch 46/100
13/13 [=====] - 0s 3ms/step - loss: 5.3051

Epoch 47/100
13/13 [=====] - 0s 3ms/step - loss: 4.8733
Epoch 48/100
13/13 [=====] - 0s 3ms/step - loss: 5.7832
Epoch 49/100
13/13 [=====] - 0s 3ms/step - loss: 4.8653
Epoch 50/100
13/13 [=====] - 0s 3ms/step - loss: 5.7771
Epoch 51/100
13/13 [=====] - 0s 3ms/step - loss: 4.5463
Epoch 52/100
13/13 [=====] - 0s 4ms/step - loss: 4.6723
Epoch 53/100
13/13 [=====] - 0s 4ms/step - loss: 4.7725
Epoch 54/100
13/13 [=====] - 0s 4ms/step - loss: 3.8493
Epoch 55/100
13/13 [=====] - 0s 5ms/step - loss: 4.4639
Epoch 56/100
13/13 [=====] - 0s 4ms/step - loss: 4.0582
Epoch 57/100
13/13 [=====] - 0s 3ms/step - loss: 3.6706
Epoch 58/100
13/13 [=====] - 0s 3ms/step - loss: 4.4159
Epoch 59/100
13/13 [=====] - 0s 3ms/step - loss: 3.9814
Epoch 60/100
13/13 [=====] - 0s 4ms/step - loss: 3.7978
Epoch 61/100
13/13 [=====] - 0s 5ms/step - loss: 3.7584
Epoch 62/100
13/13 [=====] - 0s 4ms/step - loss: 4.7628
Epoch 63/100
13/13 [=====] - 0s 4ms/step - loss: 3.7640
Epoch 64/100
13/13 [=====] - 0s 4ms/step - loss: 4.2632
Epoch 65/100
13/13 [=====] - 0s 3ms/step - loss: 3.6611
Epoch 66/100
13/13 [=====] - 0s 4ms/step - loss: 3.4057
Epoch 67/100
13/13 [=====] - 0s 3ms/step - loss: 4.2729
Epoch 68/100
13/13 [=====] - 0s 4ms/step - loss: 4.0470
Epoch 69/100
13/13 [=====] - 0s 4ms/step - loss: 4.6172
Epoch 70/100
13/13 [=====] - 0s 3ms/step - loss: 5.8933

Epoch 71/100
13/13 [=====] - 0s 4ms/step - loss: 4.0761
Epoch 72/100
13/13 [=====] - 0s 3ms/step - loss: 3.6938
Epoch 73/100
13/13 [=====] - 0s 3ms/step - loss: 3.5321
Epoch 74/100
13/13 [=====] - 0s 3ms/step - loss: 5.2945
Epoch 75/100
13/13 [=====] - 0s 3ms/step - loss: 3.7234
Epoch 76/100
13/13 [=====] - 0s 4ms/step - loss: 3.1456
Epoch 77/100
13/13 [=====] - 0s 4ms/step - loss: 3.1319
Epoch 78/100
13/13 [=====] - 0s 4ms/step - loss: 3.7920
Epoch 79/100
13/13 [=====] - 0s 3ms/step - loss: 3.1307
Epoch 80/100
13/13 [=====] - 0s 3ms/step - loss: 4.2429
Epoch 81/100
13/13 [=====] - 0s 3ms/step - loss: 3.5938
Epoch 82/100
13/13 [=====] - 0s 3ms/step - loss: 3.0202
Epoch 83/100
13/13 [=====] - 0s 3ms/step - loss: 3.1307
Epoch 84/100
13/13 [=====] - 0s 4ms/step - loss: 3.3729
Epoch 85/100
13/13 [=====] - 0s 3ms/step - loss: 2.8520
Epoch 86/100
13/13 [=====] - 0s 3ms/step - loss: 3.1072
Epoch 87/100
13/13 [=====] - 0s 3ms/step - loss: 2.5570
Epoch 88/100
13/13 [=====] - 0s 3ms/step - loss: 2.7645
Epoch 89/100
13/13 [=====] - 0s 3ms/step - loss: 2.9271
Epoch 90/100
13/13 [=====] - 0s 3ms/step - loss: 2.6950
Epoch 91/100
13/13 [=====] - 0s 3ms/step - loss: 2.8767
Epoch 92/100
13/13 [=====] - 0s 3ms/step - loss: 2.4233
Epoch 93/100
13/13 [=====] - 0s 4ms/step - loss: 2.7048
Epoch 94/100
13/13 [=====] - 0s 3ms/step - loss: 2.7012

```
Epoch 95/100
13/13 [=====] - 0s 3ms/step - loss: 2.7719
Epoch 96/100
13/13 [=====] - 0s 3ms/step - loss: 2.4201
Epoch 97/100
13/13 [=====] - 0s 3ms/step - loss: 2.4178
Epoch 98/100
13/13 [=====] - 0s 4ms/step - loss: 2.5083
Epoch 99/100
13/13 [=====] - 0s 3ms/step - loss: 2.5458
Epoch 100/100
13/13 [=====] - 0s 3ms/step - loss: 2.5559
```

```
[54]: <keras.callbacks.History at 0x7f5718059310>
```

1.2.4 Evaluando los resultados

```
[55]: y_pred = model.predict(X_test)
```

```
[56]: from sklearn.metrics import r2_score
```

```
[57]: r2 = r2_score(y_test, y_pred)
```

```
[58]: print(r2)
```

```
0.9051913696866637
```

```
[59]: from sklearn.metrics import mean_squared_error
rmse = (np.sqrt(mean_squared_error(y_test, y_pred)))
print(rmse)
```

```
2.9676084866841217
```